

## CSCI 2021 Machine Architecture and Organization, Fall 2018, Written Assignment #2

Moti Begna, begna002, sec: 007

### Problem 1 (20 points)

Consider the following assembly code for a function with a while loop:

Prob1:

```
    jmp .L2
.L5: testb $1, %dil # %dil is the lowest byte of %edi
    jne .L3
    leal 4(%rsi, %rsi, 2), %eax
    addl $5, %esi
    sall $3, %edi
    jmp .L2
.L3: leal 7(,%rdi,8), %eax
    addl $9, %esi
    sall $2, %edi
.L2: cmpl %esi, %edi
    ja .L5
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code. You may only use the source-level C variable names such as n, m, and result. Don't use register names!

```
unsigned probl(unsigned n, unsigned m) {
    unsigned result;
    while(n > m) {
        if(n == 1) {
            result = result*3 + 4;
            m = m + 5;
            n = n<<3;
        }
        else {
            result = result*8 + 7;
            m = m + 9;
            n = n << 2;
        }
    }
    return result;
}
```

## Problem 2 (20 points)

The following questions are based on the **C code on the next page**.

```
14 do_something2:
15     movl    $19, -24(%rsp)
16     movw    $-13, -8(%rsp)
17     movb    $63, -5(%rsp)
18     movq    -64(%rsp), %rax
19     movq    %rax, (%rdi)
20     movq    -56(%rsp), %rax
21     movq    %rax, 8(%rdi)
22     movq    -48(%rsp), %rax
23     movq    %rax, 16(%rdi)
24     movq    -40(%rsp), %rax
25     movq    %rax, 24(%rdi)
26     movq    -32(%rsp), %rax
27     movq    %rax, 32(%rdi)
28     movq    -24(%rsp), %rax
29     movq    %rax, 40(%rdi)
30     movq    -16(%rsp), %rax
31     movq    %rax, 48(%rdi)
32     movq    -8(%rsp), %rax
33     movq    %rax, 56(%rdi)
34     ret
```

- A. What is the output of the two print statements in lines 41 and 44? Explain this result.

Output:

do\_something1 result: 0 0 0

do\_something2 result: 19 13 ?

Explanation:

In do\_something2, the actual addresses of the bar variables are being passed through, and the values at those addresses are being changed. Thus, when b.x, b.s, and b.b are being called to print, the values that are printed are the value that were changed in those memory addresses.

In do\_something1, the values associated with b.x, b.s, and b.b were only changed locally within the function. Thus, when they are being called to print in main, the values have been unchanged and what is printed is what they were originally set to in main.

- B. In Part A, if the expected result differed from the actual result, how might you change the existing code to work as intended?

In order to print the values in do\_something2, we can change the function so that it resembles do\_something1 and pass in the reference of bar when calling the function.

C. Show the partitioning of the `bar_t` struct (i.e. show the number of bytes dedicated to each field, as well as any gaps inserted between fields).

`foo_stuffs[5] = 40`

`x = 4 bytes`

`4 bytes in gap`

`f = 8`

`s = 2`

`a = 1`

`b = 1`

`4 bytes in gap`

`64 bytes in total`

D. How many bytes does the `foo_t` union require in total?

Because `short*x` requires the most bytes out of the other variables, `foo_t` requires 88 bytes (40 for `foo_stuffs[5]` and 48 for `z[6]`).

E. Using the assembly code to the left, compiled by GCC, draw the stack frame for `do_something2`. You must show each member field for any composite structures.

$\alpha$	Caller
$\alpha + 4$	Return address
$\alpha + 8$	<code>bar_t bar</code>
$\alpha + 12$	<code>bar.x = 19</code>
$\alpha + 14$	<code>bar.s = -13</code>
$\alpha + 18$	<code>bar.b = '!</code>
$\alpha + 26$	<code>*pnt = bar</code>

```
5 typedef union foo {
6     short* x;
7     int y;
8     char z[6];
9 } foo_t;
10
11 typedef struct bar {
12     foo_t foo_stuffs[5];
13     int x;
14     double f;
15     short s;
16     char a;
17     char b;
18 } bar_t;
19
20 void do_something1(bar_t a_bar) {
21     a_bar.x = 13;
22     a_bar.s = -7;
23     a_bar.b = '!';
24 }
25
26 void do_something2(bar_t* pnt) {
27     bar_t bar;
28     bar.x = 19;
29     bar.s = -13;
30     bar.b = '?';
31     *pnt = bar;
32 }
33
34 void main() {
35     bar_t b;
36     b.x = 0;
37     b.s = 0;
38     b.b = '0';
39
40     do_something1(b);
41     printf("do_something1 result: %d %d %c\n", b.x, b.s, b.b);
42
43     do_something2(&b);
44     printf("do_something2 result: %d %d %c\n", b.x, b.s, b.b);
45 }
```

### Problem 3 (20 points)

For a C function prob3 with the general structure shown later, gcc generates the following assembly code, including a jump table:

```
prob3:
    cmpq $8, %rdx
    ja .L2
    jmp*.L4(,%rdx,8)
.L4:
    .quad .L2
    .quad .L3
    .quad .L5
    .quad .L2
    .quad .L2
    .quad .L6
    .quad .L5
    .quad .L2
    .quad .L7
.L3: leaq (%rsi, %rsi, 2), %rax
    leaq (%rax, %rax), %rsi
    addq (%rdi), %rsi
    jmp .L8
.L5: leaq (%rsi, %rsi, 2), %rax
    movq %rdx, %rax
    salq $6, %rax
    addq %rax, %rsi
    jmp .L8
.L6: leaq 80(%rsi), %rax
    movq %rax, (%rdi)
.L7: movq (%rdi), %rax
    leaq (%rax, %rsi, 4), %rsi
    jmp .L8
.L2: addq $11, %rsi
.L8: movq %rsi, (%rdi)
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code.  
You may only use the source-level C variables x, m, result, and value: don't use register names!

```
void prob3(long* value, long x, long m) {
    long result;
    switch(m) {
    case 1:
        result = value + result*6;
        break;
    case 2 :
    case 6 :
        result += (result*3)<<6;
        break;
    case 5:
        *value = 80 + result;
    case 8:
        result = value + result*4 ;
        break;
    default:
        result += 11;
    }
    *value = result;
}
```

#### Problem 4 (20 points)

```
5 fun_times:
6     pushq    %rbp
7     movq     %rsp, %rbp
8     subq     $32, %rsp
9     movq     %rdi, -24(%rbp)
10    movl     %esi, -28(%rbp)
11    cmpl     $0, -28(%rbp)
12    jne .L2
13    movl     $0, %eax
14    jmp .L3
15 .L2:
16    movl     $0, -4(%rbp)
17    movl     $0, -8(%rbp)
18    jmp .L4
19 .L5:
20    movq     -24(%rbp), %rax
21    movl     %eax, %edx
22    movl     -4(%rbp), %eax
23    addl     %eax, %edx
24    movl     -8(%rbp), %eax
25    addl     %edx, %eax
26    movl     %eax, -8(%rbp)
27    addl     $1, -4(%rbp)
28 .L4:
29    movl     -4(%rbp), %eax
30    cmpl     -28(%rbp), %eax
31    jl .L5
32    movl     -28(%rbp), %eax
33    leal     -1(%rax), %edx
34    movq     -24(%rbp), %rax
35    movl     %edx, %esi
36    movq     %rax, %rdi
37    call     fun_times
38    movl     %eax, %edx
39    movl     -8(%rbp), %eax
40    addl     %edx, %eax
41 .L3:
42    leave
43    ret
```

Using the assembly code to the left, answer the following questions.

A. What does fun\_times do? Demonstrate the logic either in C code or in pseudocode.

```
def fun_times(int sum, int i)
    if (i != 0):
        if(i > 0):
            sum += i
            i++
        fun_times(i - 1, sum)
    else:
        return sum
```

B. How many bytes are allocated on the stack with each call to fun\_times?

3 bytes

C. Where is the accumulator (i.e. the value calculated by the loop) stored?

The accumulator sum is stored in L3



### Problem 5 (20 points)

```

#define SIZE 10
void prob5(int mat[SIZE][SIZE]) {
    int r, c;
    mat[0][0] = 1;
    for(r = 1; r < SIZE; r++) {
        mat[r][0] = 1;

        for(c = 1; c < r; c++) {

            mat[r][c] = mat[r-1][c] + mat[r-1][c-1];

        }

        mat[r][r] = 1;
    }
}

```

```

1 prob5:
2     movl    $1, (%rdi)
3     movl    $1, 40(%rdi)
4     leaq    80(%rdi), %r8
5     leaq    44(%rdi), %r9
6     movl    $4, %esi
7     movl    $1, %edi
8     jmp     .L2
9 .L5:
10    movq    %r8, %rcx
11    movl    $1, (%r8)
12    cmpl    $1, %edi
13    jle     .L3
14    movl    $0, %eax
15 .L4:
16    movl    -40(%rcx,%rax), %edx
17    addl    -36(%rcx,%rax), %edx
18    movl    %edx, 4(%rcx,%rax)
19    addq    $4, %rax
20    cmpq    %rsi, %rax
21    jne     .L4
22 .L3:
23    addq    $40, %r8
24    addq    $44, %r9
25    addq    $4, %rsi
26 .L2:
27    movl    $1, (%r9)
28    addl    $1, %edi
29    cmpl    $10, %edi
30    jne     .L5
31    rep ret

```

Because the compiler has optimized some of the accesses to the array, the registers don't all correspond exactly to variables in the source code. (And the statements and instructions don't line up exactly one-to-one either, so don't put too much significance in the way we've spaced the lines). For each of the following registers, as it is used in a particular range of instructions (shown by their assembly code line number), write a C expression that corresponds to the value in the register. Your expressions should be written using the C variables `mat`, `r`, and `c`, together with C operators and constants; don't use register names.

Register	C expression
%edi, lines 10-30	mat[0][0] = 1, mat[r][0] = 1
%r8, lines 10-21	mat[r-1][0]
%eax, lines 14-21	c++
%r9, lines 24-27	mat[r][0]
%esi, lines 10-25	r++