

# CSci 5106: Programming Languages

## Procedures and Procedural Abstraction

Gopalan Nadathur

Department of Computer Science and Engineering  
University of Minnesota

Lectures in Fall 2019

## Procedures and Procedural Abstraction

Mechanism for naming a parameterized block of code

The main components of the feature

- Procedure definition

An example in C-like syntax

```
int square(int x) { return x*x; }
```

Components of the abstract syntax: *name*, *return type*,  
*formal parameters and their types*, *body*

- Procedure call

An example in C-like syntax: `square(5)`

Abstract syntax components: *name*, *actual parameters*

Syntactic wellformedness actually requires a matching between  
declared type and use

## Procedures and Restructuring the Machine

Procedures impact along both *organizational* and the *execution*  
dimensions

- The effects of procedural abstraction on the organizational  
model
  - Supports stepwise development based on problem  
decomposition
  - Basis for separating functionality from implementation  
details using the principle of *information hiding*
  - Provides the basis for realizing modularity
- The effects to the execution model
  - Allows the granularity of atomic steps to be controlled by  
the user
  - Provides for *recursion* that is a distinct way of thinking  
about computations

## Issues to Study Concerning Procedures

The main issues relate to the interaction between procedures  
and the rest of the program

- What are the mechanisms for parameter passing?
  - *Call-by-value*
  - *Call-by-reference*
  - *Call-by-value-result*
  - *Call-by-name*
- How is the name space managed?  
In particular, given a procedure or variable name
  - Which declaration governs it? Scope Issue
  - Which cell does it finally refer to? Activation Issue
- Given schemes for the above, how are they implemented?

## Parameter Passing Mechanisms

The main question is that of how actual and formal parameters match up during invocation

The standard mechanisms

- Call-by-Value (CbV)

Here we pass only the *value* of the actual parameter

- actuals can be arbitrary expressions
- if they are l-values, then they are unaffected by call

- Call-by-reference (CbR)

Here the formal parameter becomes an alias for the actual parameter

- Actual must be an l-value for this to work
- Actual can be changed by the call
- Not directly supported in C, but can be simulated by pointers and taking addresses of variables

## Call-by-Value vs Call-by-Reference

Consider the following program written in pseudo-Pascal

```
var i, j : integer;
procedure swap(x, y : integer);
  var temp : integer;
  begin temp := x; x := y; y := temp;
end;
begin i := 5; j := 7; swap(i, j); end.
```

What is the behaviour under the two parameter passing modes?

In C, we would have to write

```
void swap(int *xp, int *yp) { ... }
```

and the call would be `swap(&i, &j)`

## Call-by-Value Result

Like CbV in that actuals and formals are kept separated in a procedure invocation

However, at the end, the values of the formals are copied back into the actuals (which must be l-values)

Deviates from CbR in the presence of *aliasing*

```
var i, j : integer;
procedure foo(x, y : integer);
  begin i := y; end;
begin i := 2; j := 3; foo(i, j); end
```

Consider what happens in the two cases

Was introduced in *Ada* to draw attention to aliasing and program behaviour, but is not a commonly used mechanism

## Names, Binding and Bound Occurrences

- Names play an important role in programs

They identify variables, procedures, constants, types, parameters, etc

- Corresponding to names, there are two important concepts

- declarations introducing them with their properties
- name uses

These two kinds of uses are called *binding* and *bound* occurrences

- It is important to be able to say which binding occurrence pertains to which bound occurrence
- Scope rules provide the mechanisms for doing this

## Scoping Paradigms

These come in two flavours

- *Static or lexical scoping*

*Guiding Principle:* Binding occurrence should be determinable purely from the text

- Dynamic scoping

Binding occurrences depend on execution pattern

In particular, procedure invocations “activate” binding occurrences that govern subsequent usage

Note: a distinction arises *only when* there are non-local name occurrences

## An Example that Differentiates the Two Notions

```
program testscope;
var i,j : integer;

procedure W;
begin i := j; end;

procedure D;
var j : integer;
begin j := 5; W; end;

begin j := 3; D; write(i); end.
```

What value is printed in each case?

## Which Scoping Paradigm is Better?

The answer to this question is related to issues of modularity and abstraction

*We should be able to understand what a block of code does independently of where it is going to be used*

Lexical scoping supports this principle

*If a procedure affects a nonlocal variable, the variable that is changed is known from the textual context in which the procedure occurs*

Dynamic scoping *does not* support this principle

For this reason, dynamic scoping is seldom used

It has actually been a red herring: McCarthy “introduced” it in Lisp but long ago said it was a bug

## Call by Name

This is a parameter passing mechanism that is based on macro expansion done right

Here is the way to understand how it works

- Textually substitute actuals for formals in the procedure body
- Textually substitute resulting code at places of call

However, when doing the substitution, *rename* variables to make sure to avoid accidental capture

Such renaming is justified: particular names should not be relevant to the meanings of programs

## Example of Call by Name

```
int i, j;
void bar1 (int x) {
    int j;
    i = x; }
void bar2() {
    int i;
    bar1(j); }
int main() {
    ... bar2(); ... }
```

Now consider what happens to the call `bar1(j)`

To get it right we must

- rename the `j` in `bar1` when replacing `x` by `j`
- rename the `i` in `bar2` when replacing `bar1(j)`

## Call-by-Name versus Call-by-Reference

CbN is like CbR but with *lazy address calculations*

Delayed address calculation can make a difference

For example, consider the following code

```
procedure swap(x,y : integer);
var temp : integer;
begin
    temp := x; x := y; y := temp;
end;

begin
    i := 1; A[1] := 2; A[2] := 2;
    swap(i,A[i]);
end.
```

Call-by-reference: `i = 2; A[1] = 1; A[2] = 2;`

Call-by-name: `i = 2; A[1] = 2; A[2] = 1;`

## Nested Scopes and Visibility

Binding occurrences have scopes over blocks of text

When such blocks can be nested, there is a possibility for creating visibility “holes”

```
{
    int i;
    . . .
    {
        char i;
        . . .
    }
    . . .
}
```

The inner declaration of `i` creates an *occlusion* zone for the outer declaration

The outer declaration is still active but just not visible

## Procedure Activation

Refers to the commencement of execution of a procedure

The questions that are relevant to this

- How do we deal with space for local variables and parameters?
  - where is this space to be provided?
  - how do we locate the cell for a particular variable?
- What happens after the invocation completes?
  - how do we reclaim the space that was allocated?
  - how do we return the results?
  - Where should control return to?

The main complexity arises from the presence of recursion

## Recursion and Procedure Activation

Recursion makes it impossible to allocate space and to identify the reference of names statically

```
void ReverseLine() {  
    char ch;  
    while (getc(ch) != '\n') {  
        ReverseLine(); write(ch);  
    }  
}
```

The number of incarnations of `ReverseLine`, and, hence, of `ch` is *a priori* unknown

In this situation we have to think of a *dynamic* scheme for dealing with both issues

## Stack Based Activation Records

The common procedure activation model makes the following assumptions

- Procedures complete their work in reverse order of invocation
- Procedure space is not needed after completion; specifically, local variables are not needed

The assumptions do not always hold, e.g. in functional programming and for co-routines

When they hold, space can be provided and gathered back in a stack-based fashion

There is often hardware support for maintaining such stacks, known as *activation record* stacks

## Activation Record Structure

Activation records contain information of the following sort:

- Space for parameters, local variables, etc
- Return address
- Pointer to the activation record of the runtime caller  
*Dynamic or Control Link*
- Pointer to the activation record of the nearest statically enclosing procedure (for lexical scoping)  
*Static or Access Link*

The architecture designer usually describes a standard format to enhance language interoperability

## Resolving Variable References

The issue here is that of finding the cell in an activation record for a variable use

Usefully thought of as a two-step process:

- Figure out the connection between variable use and binding occurrences  
*Scope issue, resolved at compile-time*
- Figure out where the space is for the binding occurrence for the relevant invocation  
*Run-time connection between the procedure and its activation record*

Once we have understood the latter issue, then we can design a composition of the two mappings

## Activation Records and Variable Reference (Example)

We can illustrate these issues by considering the following code

```
program M;
  var b, c;
  procedure P;
    var x, y, z;
    procedure Q;
      var b;
      begin (* Q *) b = x; ... R ... end;
    procedure R;
      var c;
      begin (* R *) x = b + c; ... P ... end;
    begin (* P *) ... Q ... end;
  begin (* M *) ... P ... end.
```

Consider now how activation records are set up, how access links are calculated and how variable references are resolved

## Setting Up Access Links

One way to do this is for the caller to pass the access link as a parameter in a procedure call

The access link can be calculated using the *nesting depth* of a procedure, i.e. the “level” at which a procedure is declared

More concretely, the compiler can set up code to pass the static link parameter as follows

- Statically determine the difference between the nesting depths of the caller and the callee
- Generate code to chain back through a number of access links equal to one more than this difference
- Insert this code at the place where the access link needs to be passed as a parameter

## Resolving Variable References

Once again, the nesting depths are useful for this

Associated with a variable declaration is a nesting depth and an *offset*

Associated with the variable use is also a nesting depth

The compiler can then generate access code that carries out the following steps:

- Chain back through a number of activation records equal to the difference between declaration and use nesting depths
- Use offset associated with the declaration to get to the right cell

## Speeding Up Variable Access

Chaining back through activations records at runtime can be avoided using the following ideas

- Maintain a vector whose  $i^{th}$  entry points to the activation record at nesting level  $i$ , if it is active  
This vector is called a *display*
- To look up a variable declared at nesting depth  $i$ 
  - get the pointer to the activation record from  $i$ th display entry
  - use the offset information for the variable to determine the cell to access
- At a call to a procedure at nesting depth  $i$ 
  - Save value in the  $i$ th display cell in the activation record
  - Save a pointer to the new activation record in the  $i$ th cell
  - Restore the  $i$ th cell in display upon return from the call

## Procedures as Parameters

The following code shows the kind of scenario that is of interest

```
procedure P;  
  var x;  
  procedure Q;  
    begin ... x ... end;  
  procedure R(proc X);  
    begin ... X ... end;  
  begin ... R(Q) ... R(P) ... end
```

The question: How do we set the access link for the call to `x` from within `R`?

One solution: Pass the access link along as one component of the procedure parameter